

# VISMAY JAIN

+91-6351464527 | [ai20.vismay.jain@gmail.com](mailto:ai20.vismay.jain@gmail.com) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## PROFESSIONAL SUMMARY

GenAI Engineer specializing in enterprise-grade agentic systems and LLM-powered automation. Proven track record in architecting LangGraph-based agent frameworks, fine-tuning large language models, and building production-ready RAG systems. Expert in designing full-stack AI solutions with observability and cost optimization.

## EDUCATION

Uka Tarsadia University

Bachelor of Technology in Artificial Intelligence and Data Science | CGPA: 8.7/10

Bardoli, Gujarat

Oct 2020 – June 2024

## TECHNICAL SKILLS

**Programming Languages:** Python, SQL, TypeScript, C#

**AI/ML Frameworks:** LangGraph, LangChain, FastAPI, Hugging Face Transformers, vLLM, Ollama

**LLMs & GenAI:** Claude, LLaMA, GPT, PEFT, QLoRA, RAG Architecture, Prompt Engineering, MCP

**ML & Clustering:** HDBSCAN, UMAP, scikit-learn, NumPy, embeddings, dimensionality reduction

**Web Frameworks:** FastAPI, Flask, React.js, Node.js, REST APIs, Server-Sent Events (SSE)

**DevOps & Infrastructure:** Docker, AWS (S3, Lambda, EC2), GitHub Actions, Terraform, CI/CD

**Databases:** MongoDB, Redis, SQLite, PostgreSQL, Pinecone, ChromaDB, Vector Databases

**Observability & Monitoring:** Langfuse, Prometheus, trace logging, SLA monitoring, token tracking

**Integration Tools:** Git, Postman, Zoom SDK, Slack APIs, Whisper, MediaPipe, JSON-RPC

## PROFESSIONAL EXPERIENCE

GenAI Engineer

Ideyalabs

June 2025 – Present

Remote

- Architected **Dev Agent** – multi-agent code generation platform transforming user stories into production-ready full-stack code using LangGraph orchestration with 5 specialized agents (Architecture, Screen Planner, Frontend, Backend, Database)
- Built **QA Agent** – AI-powered test automation platform with document intelligence aggregating multi-source docs, LLM-driven semantic understanding, and hybrid vector search using Pinecone, achieving **95%+ test coverage**
- Architected **CSX AI Customer Support Agent** serving 10,000+ conversations monthly with 99.9% uptime, reducing response latency by **90%** (110s → 11s) through intelligent caching, async processing, and real-time SSE streaming
- Built custom **Model Context Protocol (MCP)** client achieving **98% faster tool execution** (100s → 2s) with JSON-RPC 2.0 protocol support, automatic tool discovery, and autonomous agentic loop
- Designed LangGraph workflow orchestration system with 15+ reusable nodes and 6 smart routers, enabling **70% faster workflow development** with checkpoint-based state persistence and real-time event streaming
- Optimized performance with asyncio semaphores for dynamic parallel processing reducing generation time by **50%**, with multi-provider LLM fallback (OpenAI → Anthropic) and comprehensive validation framework
- Integrated **Langfuse** observability platform enabling prompt versioning, distributed trace logs, token usage analytics, and SLA monitoring across agent workflows
- Architected **AutoPrompt JIRA Integration** – full-stack AI prompt generation system deployed as Atlassian Forge Custom UI Panel, serving 100+ concurrent users with user-segregated PostgreSQL architecture and CASCADE relationships
- Optimized parallel LLM execution by bypassing request queues with direct Anthropic API calls using `asyncio.gather()`, reducing generation time by **54%** (41s → 18s) while delivering 3 prompt variations within Forge's 25-second timeout constraint
- Built **DocAgent** – autonomous documentation platform using dedicated LLM agents (Epic, Feature, Interlocking) grounded via Knowledge Graph traversal + Pinecone RAG, with SSE streaming and React Flow visualization, reducing documentation time from **days to minutes**

AI Engineer

Commercient LLC

June 2024 – June 2025

Remote

- Architected and deployed in-house LLM-powered sales assistant replacing third-party DriftBot, achieving **\$12,000 annual cost savings** while improving response accuracy by 35% through custom fine-tuning

- Engineered multi-modal chatbot processing documents, video, and images using Whisper ASR, LLaMA Vision, and GPT Vision API with integrated Pinecone RAG pipeline
- Deployed fine-tuned LLMs using **vLLM** inference engine over multi-GPU nodes (A100), reducing inference latency by **40%** while maintaining 99.5% uptime
- Built automated helpdesk ticket summarization pipeline processing **200K+ tokens** using LangChain, reducing manual triage time by 75% and improving CSAT scores by 28%

#### AI/ML Intern

Dec 2023 – June 2024

*Commercient LLC*

*Remote*

- Built production RAG system using Pinecone vector database and OpenAI embeddings, reducing response time from 20s to 5s through semantic caching and query optimization
- Designed LangChain-based helpdesk automation processing 200K tokens for ticket classification and intelligent routing, achieving **85% automation rate**

#### AI/ML Intern

May 2023 – July 2023

*Logictrix Infotech*

*Surat, India*

- Fine-tuned ALBERT model for enterprise email classification achieving **92% accuracy**, deployed on AWS Lambda with automatic scaling

## KEY PROJECTS

---

### Vyxion – Enterprise Knowledge Graph & AI Scenario Intelligence Platform

2024 – Present

- Architected full-stack decision intelligence platform using FastAPI, Next.js 16, MongoDB, and Claude AI; built custom graph engine with NetworkX for complex relationship modeling and multi-timeline scenario simulation (optimistic/realistic/pessimistic) with probabilistic outcome modeling
- Engineered interactive visualization handling **500+ nodes at 60fps** via vis-network, natural language scenario generation, Monte Carlo simulations (1000+ iterations), bias detection across 6 cognitive bias types, and audit trail versioning – achieving **sub-500ms** graph response times

### RAG-based Code Documentation LLM Agent

2024

- Designed microservice for Retrieval-Augmented Generation adapting code documentation websites into intelligent LLM agents, processing **10K+ pages**
- Employed Pinecone vector store with hybrid search and metadata filtering, ensuring sub-second query resolution with **95% answer accuracy**

### Sierra: Voice-Activated GenAI Assistant

2024

- Created voice interface using Python + Groq API + LLaMA for intent parsing and natural language generation, integrating Spotify, WhatsApp, and DALL-E APIs

## CERTIFICATIONS

---

- **AWS Certified Machine Learning – Specialty** – Amazon Web Services
- **AWS Partner Accreditation – Technical** – Amazon Web Services
- **IBM Data Analysis Professional Certificate** – IBM